

Data Profile & Processing

Complaint → Regulation Classifier (CMPL-REG-24)
Curated CFPB consumer-complaint narratives

Generated: 2026-07-22 19:37 UTC
Source: CFPB Consumer Complaint Database (public)
Dataset: data/complaints/cfpb_complaints.csv - 4,000 rows
Model: CMPL-REG-24 - repo: reg-agents

1 · Processing pipeline

stage	what it does	parameter / evidence
1 · Acquire	Download from the public CFPB Cor	scripts/fetch_cfpb_complaints.py parse of the bu
2 · Length filter	Keep narratives ≥ 120 chars; truncat	MIN_CHARS=120 · MAX_CHARS=1800
3 · Exact dedup	Hash of whitespace/case-normalize	0 duplicates remain (verified below)
4 · Near dedup	First-200-chars normalized fingerpr	removes boilerplate template complaints
5 · PII check	CFPB pre-masks PII as 'XXXX'; verifi	82% of narratives carry masking tokens
6 · Balanced sampling	Per product-issue cap to fight credit	PER_ISSUE_CAP=350 · max observed 126
7 · Weak labeling	CFPB product/issue taxonomy + key	reg_agents/common/complaints.py REGULATORY
8 · Scoring holdout	5% stratified reserve carved out FIF	holdout 200 (193 / 7) h, never used in
9 · Split	80/10/10 train/val/test on the rema	train 3,040 / val 380 / test 380 ked seed

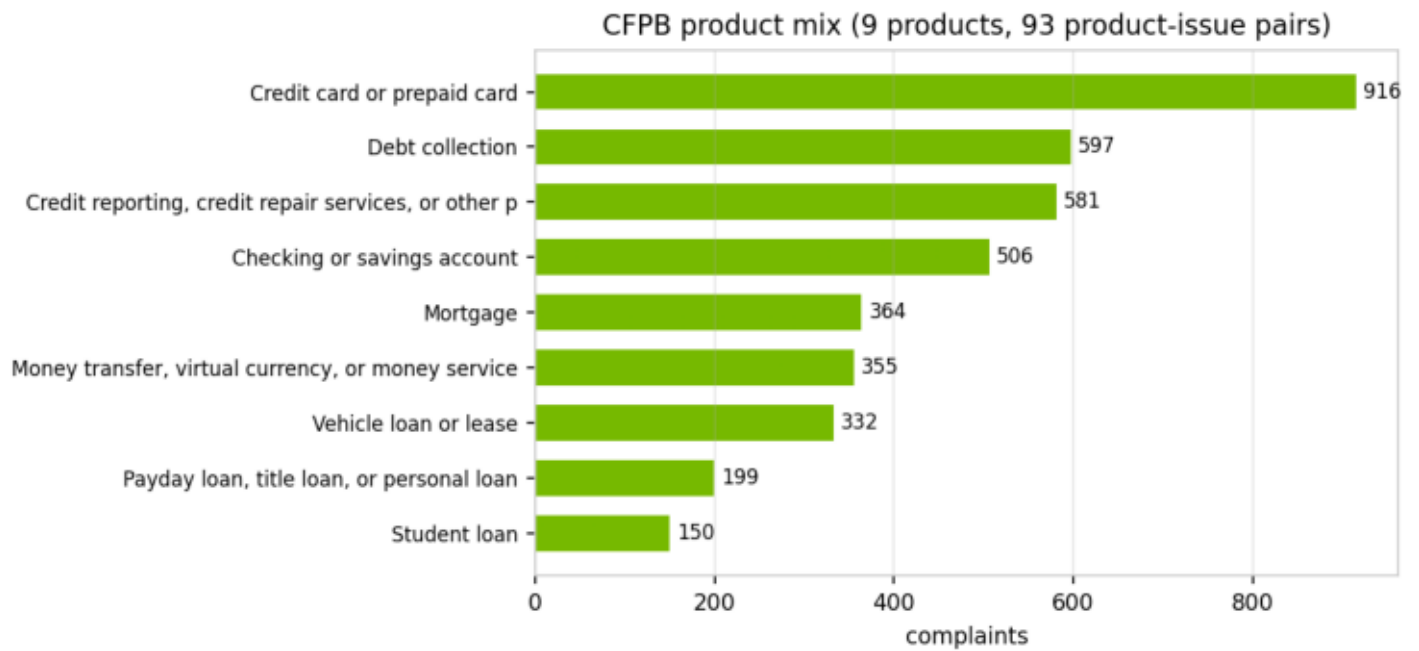
2 · Dataset schema

column	dtype	non-null	unique
complaint_id	int64	4,000	4,000
date_received	object	4,000	67
product	object	4,000	9
sub_product	object	4,000	47
issue	object	4,000	73
sub_issue	object	2,970	146
company	object	4,000	557
state	object	3,986	56
tags	object	769	3
narrative	object	4,000	4,000
label	object	4,000	24
is_regulatory	int64	4,000	2

3 · Composition

The dataset is NOT all-regulatory: 3,865 of 4,000 narratives (96.6%) carry a weak regulatory label and 135 (3.4%) are NON_REGULATORY service complaints. The imbalance is a property of the source: the CFPB database predominantly receives complaints with a regulatory nexus, and the curation deliberately keeps the natural mix rather than rebalancing, so stage-1 metrics reflect production-like prevalence. This is why PR-AUC (not accuracy or ROC-AUC) is the primary stage-1 metric, why the classifiers use class weighting (`class_weight='balanced'; scale_pos_weight`), and why the score-distribution and calibration figures in the development document matter more than headline accuracy.

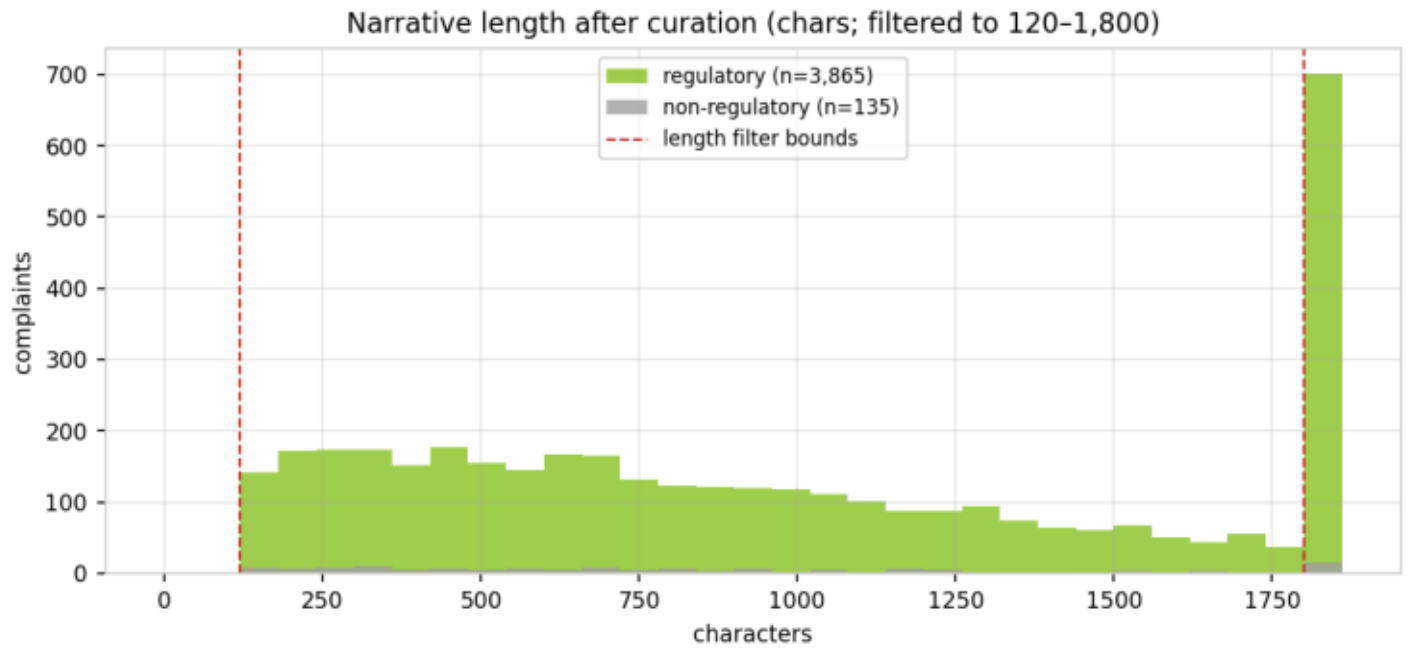
3 · Product mix



4 · Label coverage

label	category	n	share
UDAAP	UDAAP — Unfair, Deceptive	680	17.0%
LOAN_SERVICING	Loan Servicing (Auto/Student)	472	11.8%
REG_E_UNAUTHORIZED	Reg E / EFTA — Unauthorized	276	6.9%
BSA_AML	BSA/AML — Account Freeze	216	5.4%
TILA_ORIGINATION	TILA — Credit Origination &	213	5.3%
FDCPA_THREATS	FDCPA — False Statements	208	5.2%
FDCPA_DEBT_VALIDATION	FDCPA — Debt Validation /	207	5.2%
REG_E_ERROR_RESOLUTION	Reg E / EFTA — Error Resol	197	4.9%
REG_Z_DISCLOSURE	Reg Z / TILA — Fees, Inter	196	4.9%
FCRA_ACCURACY	FCRA — Accuracy of Repor	179	4.5%
RESPA_SERVICING	RESPA — Mortgage Servi	168	4.2%
REG_Z_BILLING	Reg Z / FCBA — Billing Err	156	3.9%
FDCPA_COMMUNICATION	FDCPA — Communication	138	3.5%
GLBA_PRIVACY	GLBA / Privacy — Informat	136	3.4%
NON_REGULATORY	Non-Regulatory — General	135	3.4%
RESPA_LOSS_MITIGATION	RESPA — Loss Mitigation &	109	2.7%
FCRA_INVESTIGATION	FCRA — Reinvestigation of	107	2.7%
ECOA_DISCRIMINATION	ECOA / Reg B — Credit Dis	61	1.5%
FCRA_PERMISSIBLE_PURPOSE	FCRA — Permissible Purpos	48	1.2%
TISA_REG_DD	TISA / Reg DD — Deposit A	36	0.9%
REG_CC_FUNDS	Reg CC — Funds Availabilit	26	0.7%
SCRA_MLA	SCRA / MLA — Servicemen	13	0.3%
SALES_PRACTICES	Sales Practices — Unautho	12	0.3%
ECOA_ADVERSE_ACTION	ECOA / Reg B — Adverse A	11	0.3%

5 · Narrative length



6 · Scoring holdout + train/val/test design

Before any modeling split, a stratified 5% SCORING HOLDOUT is reserved (fixed seed). It simulates a fresh batch arriving through the ingestion layer – scripts/score_batch.py and the UI's batch-scoring upload feed it through the two-stage pipeline and emit a scored CSV (complaint id, complaint, stage-1 score, LLM reasoning). No training, validation, test, or threshold-tuning step ever touches these rows.

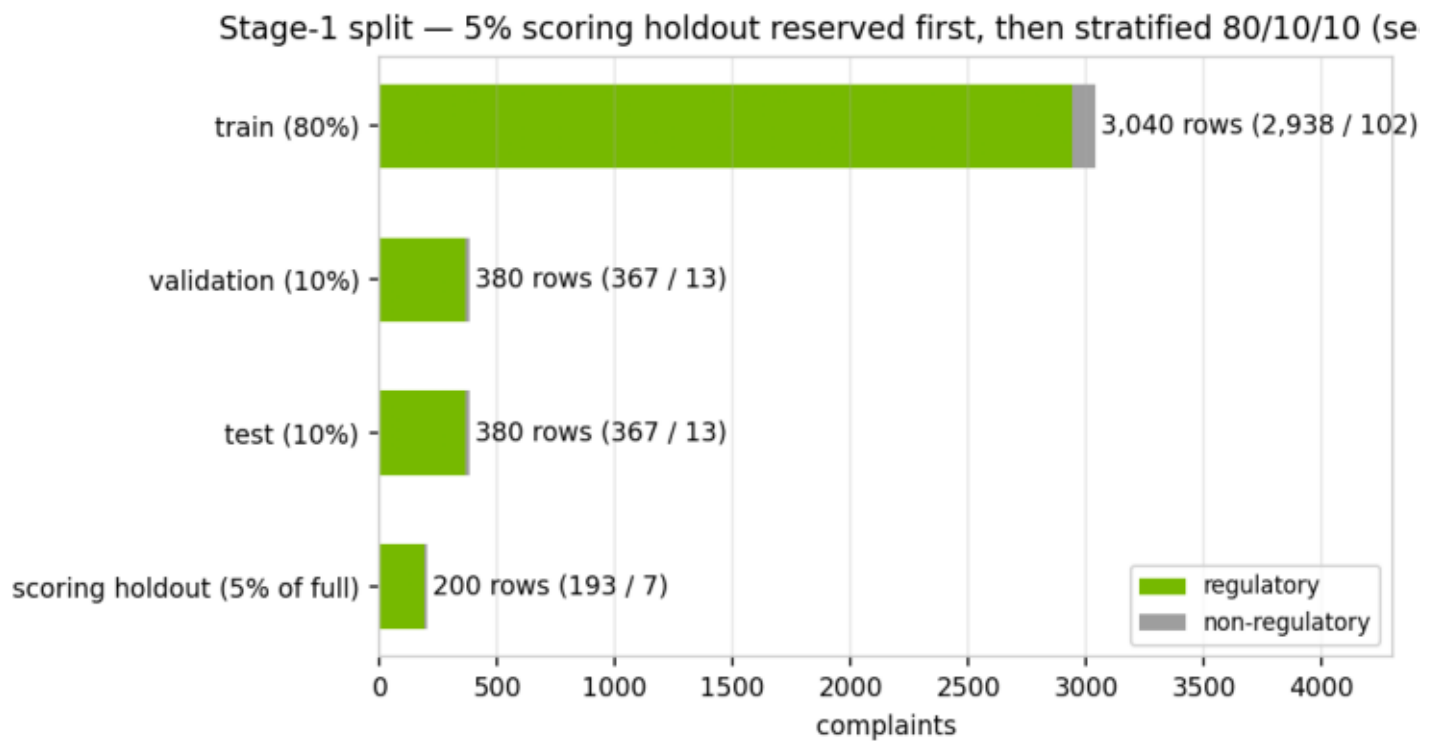
Stage 1 then uses a stratified 80/10/10 train/validation/test split of the remaining 95% (fixed seed): the test set is split off first (10%), then the remainder is split 8/1 into train and validation. Champion selection happens on validation PR-AUC only, and the decision cut-off on P(regulatory) is optimized on the same validation fold (maximizing minority-class F1) instead of assuming the default 0.5. The test set is touched exactly once, to report final metrics for the already-selected, already-thresholded model – so neither model choice nor cut-off choice can leak into reported performance.

Stage 2 involves no training: it is a prompted LLM over retrieved context. Its evaluation uses a separate stratified sample (n=115) of regulatory complaints scored against the weak labels, reported in the validation report. No complaint text is used to fit stage-2 parameters.

6 · Split composition

split	rows	regulatory	non-regulatory	reg rate
train	3,040	2,938	102	96.64%
validation	380	367	13	96.58%
test	380	367	13	96.58%
scoring holdout (5% c	200	193	7	96.50%

6 · Split (stratified)



7 · Data-quality checks (recomputed)

check	result	expectation
rows	4,000	= TARGET_ROWS (4,000)
missing values (key columns)	0	complaint_id, product, issue, narrative
exact duplicate narratives	0	post-normalization
length bounds respected	min 120 · max 1800	filter is 120–1,800 chars
PII masking present	82.3% of narratives	CFPB 'XXXX' masking tokens
max complaints per product-issue	126	cap is 350
label coverage	24 of 24 categories	all categories populated
regulatory / non-regulatory	3,865 / 135	96.6% regulatory